



# ALTO over New Transport

draft-ietf-alto-new-transport-01

Roland Schott  
Y. Richard Yang  
Kai Gao  
Jensen Zhang

July 26, 2022

IETF 114

# Outline

- Document activities
- High-level recap
- Major changes from IETF 113
- Discussions and remaining issues to finalize

# Document Activities

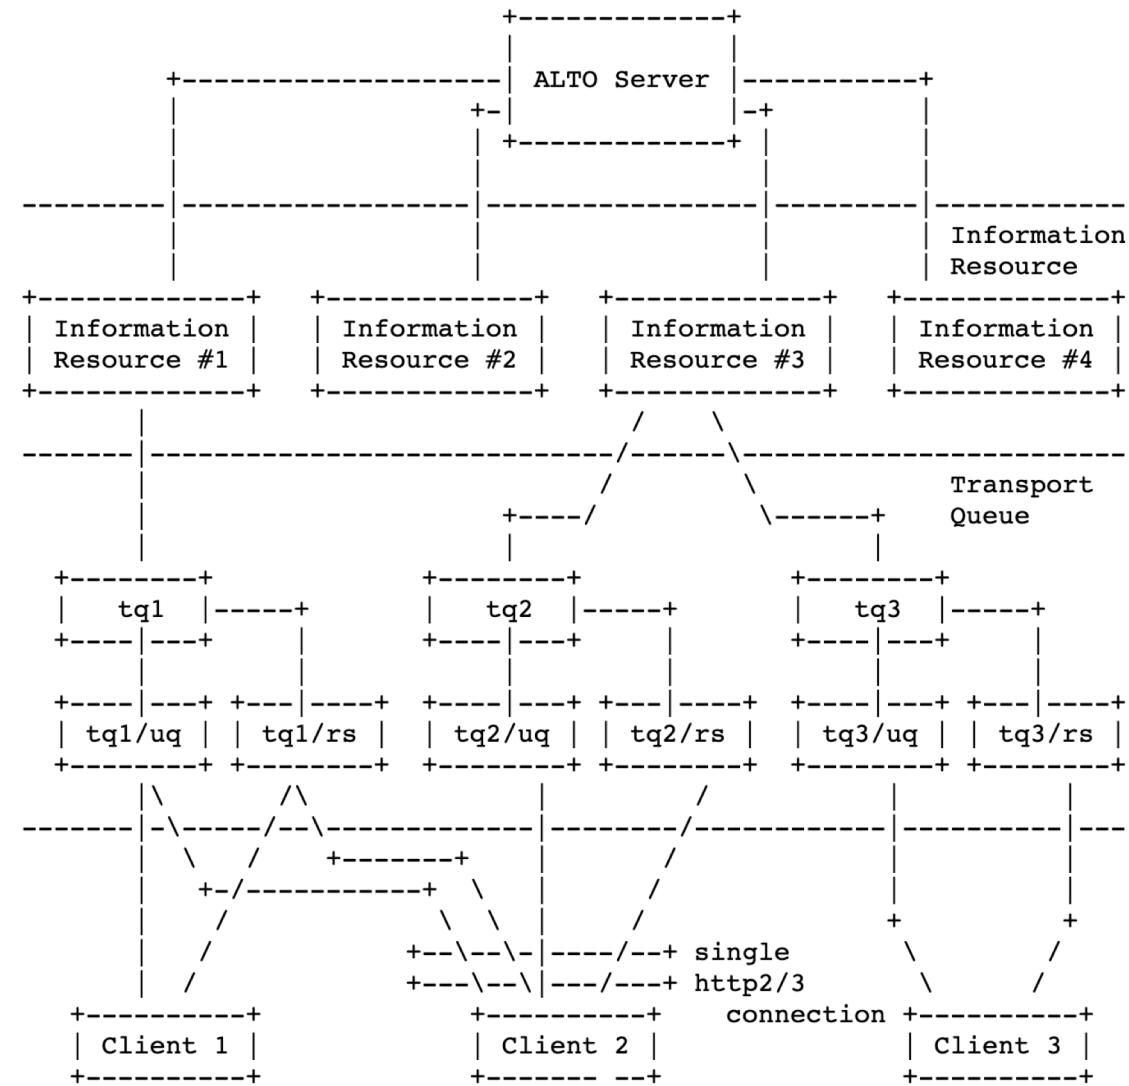
- Adopted as working group (WG) document
  - draft-ietf-alto-new-transport-00 submitted on June 22, 2022
- Updates to fix review feedbacks
  - Many helpful feedback from WG chairs (Med and Qin) and members (Luis, Sabine, Jordi)
  - draft-ietf-alto-new-transport-01 submitted on July 10, 2022
- Three excellent, early HTTP expert reviews
  - [MT] Martin Thompson (July 11)
    - [https://mailarchive.ietf.org/arch/msg/alto/sa1Pv7jmTfBF3TbGuJr\\_PffljXg/](https://mailarchive.ietf.org/arch/msg/alto/sa1Pv7jmTfBF3TbGuJr_PffljXg/)
  - [SD] Spencer Dawkins (July 15)
    - <https://datatracker.ietf.org/doc/review-ietf-alto-new-transport-01-artart-early-dawkins-2022-07-15/>
  - [MN] Mark Nottingham (July 17)
    - <https://mailarchive.ietf.org/arch/msg/alto/D84S0qLbgtpL0-jf93gNPS3NUJE/>
  - More discussion details see WG mailing list archive

# Recap: ALTO New Transport Design Requirements

- From ALTO base protocol [RFC 7285]
  - R0: Client can request any ALTO resource using the connection, just as using ALTO base protocol using HTTP/1.x
- From ALTO SSE [RFC 8895]
  - R1: Client can request the addition (start) of incremental updates to a resource
  - R2: Client can request the deletion (stop) of incremental updates to a resource
  - R3: Server can signal to the client the start or stop of incremental updates to a resource
  - R4: Server can choose the type of each incremental update encoding, as long as the type is indicated to be acceptable by the client
- From ALTO base framework [RFC 7285]
  - R5: Design follows basic HTTP Representational State Transfer architecture if possible
    - Can use only a limited number of verbs (GET, POST, PUT, DELETE, HEAD)
  - R6: Design takes advantage of HTTP/2 design features such as parallel transfer and respects HTTP/2 semantics [PUSH\_PROMISE]
- Allow flexible deployment
  - R7: Capability negotiation

# Recap: New Transport Architecture

- An ALTO server provides multiple info resources {R[i]}
  - URIs announced in Information Resource Directory (IRD)
- An ALTO client uses **a single HTTP/[2-3] connection** to receive the content of multiple resources
- Content of a resource with updates stored in an update queue (URI)
- Client can pull items in the update queue from server or server can push items in the update queue to client



# Major Changes from IETF 113

- Substantial updates to specify protocol details

1.	Introduction . . . . .	2
2.	ALTO New Transport Design Requirements . . . . .	3
3.	ALTO New Transport Information Structure . . . . .	4
3.1.	Transport Queue . . . . .	5
3.2.	Incremental Update Message Queue . . . . .	5
3.3.	Examples . . . . .	5
4.	ALTO/H2 Information Resource Directory (IRD) . . . . .	10
5.	Security Considerations . . . . .	13
6.	IANA Considerations . . . . .	13
7.	Acknowledgments . . . . .	14
8.	References . . . . .	14
8.1.	Normative References . . . . .	14
8.2.	Informative References . . . . .	14
	Authors' Addresses . . . . .	15

1.	Introduction . . . . .	3
2.	ALTO/H2 Design Requirements . . . . .	4
3.	ALTO/H2 Design Overview . . . . .	4
4.	Transport Queue . . . . .	7
4.1.	Transport Queue Operations . . . . .	7
4.2.	Examples . . . . .	8
5.	Incremental Updates Queue . . . . .	10
5.1.	Incremental Updates Queue Operations . . . . .	11
5.2.	Examples . . . . .	11
6.	Individual Updates . . . . .	12
6.1.	Individual Updates Operations . . . . .	12
6.2.	Examples . . . . .	13
7.	Receiver Set . . . . .	16
7.1.	Receiver Set Operations . . . . .	16
7.2.	Examples . . . . .	17
8.	ALTO/H2 Stream Management . . . . .	17
8.1.	Objectives . . . . .	17
8.2.	Client -> Server [Create Transport Queue] . . . . .	17
8.3.	Client -> Server [Close Transport Queue] . . . . .	18
8.4.	Client -> Server [Request on Data of a Transport Queue on Stream SID_tq] . . . . .	18
8.5.	Server -> Client [PUSH_PROMISE for Transport Queue on Stream SID_tq] . . . . .	18
8.6.	Concurrency Management . . . . .	18
9.	ALTO/H2 Information Resource Directory (IRD) . . . . .	19
10.	Security Considerations . . . . .	22
11.	IANA Considerations . . . . .	22
12.	Acknowledgments . . . . .	22
13.	References . . . . .	22
13.1.	Normative References . . . . .	22
13.2.	Informative References . . . . .	23
	Appendix A. Outlook to ALTO with HTTP/3 . . . . .	23
	Authors' Addresses . . . . .	24

# Discussions and Remaining Issues

- Excellent comments and discussions from HTTP experts
  - Main remaining issues for the WG are to finalize the **concurrency control** and **semantics** of the transport of the information items
  - They may represent generic problem for HTTP/[2-3] design
  - The finalization of the remaining issues does not look like to need substantial texts, but need careful WG decision and specification

# Discuss 1: Finalizing Op Mode: Client Pull/Server Push/ Client Long Pull/Server Put

- Current mechanisms
  - Client pull

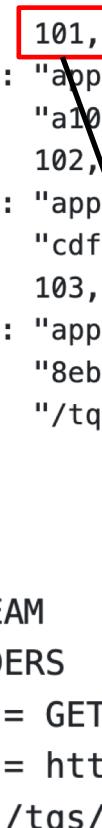
Client opens a connection to the server  
Client opens/identifies a transport queue tq

```
// pull mode
Client requests transport queue status of tq
Client requests an element in the incremental update queue
```

```
// push mode
Client becomes a receiver
Client receives incremental push updates
Client closes the transport queue tq
Client closes the connection
```

Figure 3: ALTO New Transport Information Structure.

```
{
  [
    {"seq": 101,
     "media-type": "application/alto-costmap+json",
     "tag": "a10ce8b059740b0b2e3f8eb1d4785acd42231bfe" },
    {"seq": 102,
     "media-type": "application/merge-patch+json",
     "tag": "cdf0222x59740b0b2e3f8eb1d4785acd42231bfe" },
    {"seq": 103,
     "media-type": "application/merge-patch+json",
     "tag": "8eb1d4785acd42231bfecdf0222x59740b0b2e3f",
     "link": "/tqs/2718281828459/snapshot/2e3f"}
  ],
  HEADERS
  + END_STREAM
  + END_HEADERS
  :method = GET
  :scheme = https
  :path = /tqs/2718281828459/uq/101
  host = alto.example.com
  accept = application/alto-error+json,
            application/alto-costmap+json
```



# Discuss 1: Finalizing Op Mode: Client Pull/Server Push/ Client Long Pull/Server Put

- Current mechanisms

- Client pull
- Server push (push\_promise)

```
Client opens a connection to the server  
Client opens/identifies a transport queue tq  
  // pull mode  
Client requests transport queue status of tq  
Client requests an element in the incremental update queue
```

```
// push mode  
Client becomes a receiver  
Client receives incremental push updates
```

```
Client closes the transport queue tq  
Client closes the connection
```

Figure 3: ALTO New Transport Information Structure.

Server → client PUSH\_PROMISE in current stream:

PUSH\_PROMISE

- END\_STREAM

Promised Stream 4  
HEADER BLOCK

- :method = GET
- :scheme = https
- :path = /tqs/2718281828459/uq/101
- host = alto.example.com
- accept = application/alto-error+json,  
application/alto-costmap+json

Server → client content Stream 4:

HEADERS

- + END\_STREAM
- + END\_HEADERS

- :status = 200
- content-type = application/alto-costmap+json
- content-length = TBD

# Discuss 1: Finalizing Op Mode: Client Pull/Server Push/ Client Long Pull/Server Put

- Incremental push alternatives

[MN review]

- Client long pull

- Allow request on next seq

Client opens a connection to the server

Client opens/identifies a transport queue tq

// pull mode

Client requests transport queue status of tq

Client requests an element in the incremental update queue

// push mode

Client becomes a receiver

Client receives incremental push updates

Client closes the transport queue tq

Client closes the connection

Figure 3: ALTO New Transport Information Structure.

```
{  
  [  
    {"seq": 101,  
     "media-type": "application/alto-costmap+json",  
     "tag": "a10ce8b059740b0b2e3f8eb1d4785acd42231bfe"},  
    {"seq": 102,  
     "media-type": "application/merge-patch+json",  
     "tag": "cdf0222x59740b0b2e3f8eb1d4785acd42231bfe"},  
    {"seq": 103, 103,  
     "media-type": "application/merge-patch+json",  
     "tag": "8eb1d4785acd42231bfecdf0222x59740b0b2e3f",  
     "link": "/tqs/2718281828459/snapshot/2e3f"}  
  ],  
  HEADERS  
  + END_STREAM  
  + END_HEADERS  
  :method = GET  
  :scheme = https  
  :path = /tqs/2718281828459/uq/ 104  
  host = alto.example.com  
  accept = application/alto-error+json,  
           application/alto-costmap+json
```

# Discuss 1: Finalizing Op Mode: Client Pull/Server Push/ Client Long Pull/Server Put

- Incremental push alternatives

- Client long pull
  - **Server put [MN review]**

- Benefit:

- Avoid “awkward” promise  
(current spec: MUST NOT cancel push promise)

- Issue: semantics

- Client conceptually is only a cache, not a persistent state replica

Server → client PUSH\_PROMISE in current stream:

~~PUSH\_PROMISE~~

- END\_STREAM
- ~~Promised Stream 4~~
- HEADER BLOCK
  - :method = ~~GET~~ PUT
  - :scheme = https
  - :path = /tqs/2718281828459/uq/101
  - host = alto.example.com
  - accept = application/alto-error+json,  
application/alto-costmap+json

Server → client content Stream 4:

HEADERS

- + END\_STREAM
- + END\_HEADERS
- :status = 200
- content-type = application/alto-costmap+json
- content-length = TBD

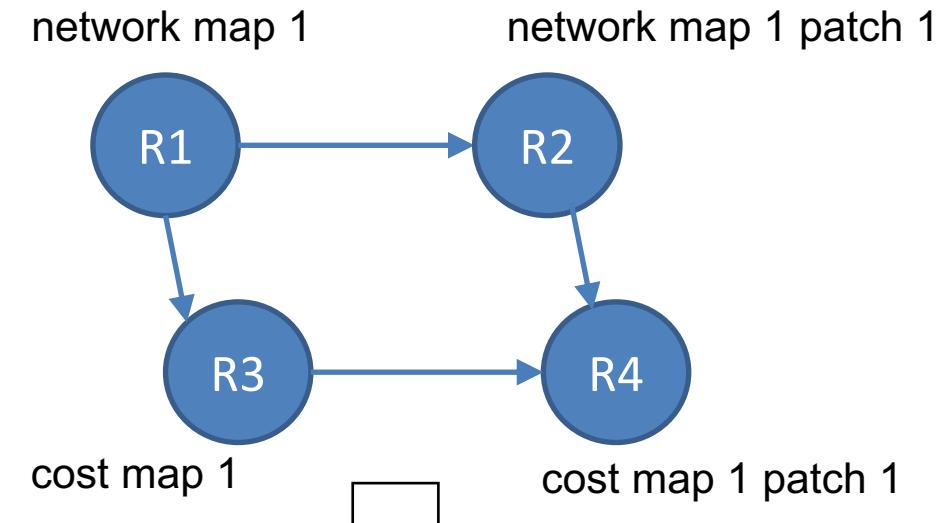
## Discuss 2: How Much to Specify

### Ordering Control: Transport-Aware of App Semantics [MN, MT, SD reviews]

App/ALTO

Design 1:

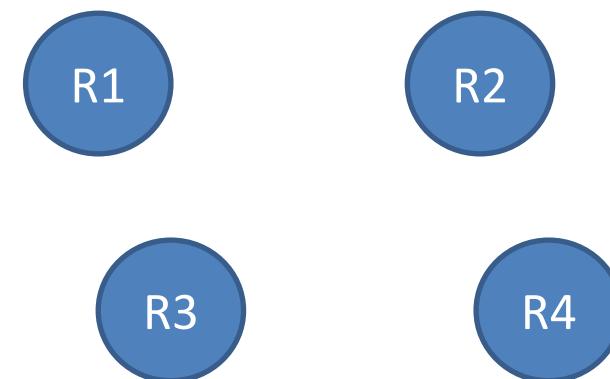
- ALTO specifies only: mapping each  $R_i$  to an independent HTTP/2-3 stream; HTTP does scheduling.
- Issue: HTTP could schedule  $R_4$ , then  $R_3$ , then  $R_2$  and then  $R_1$  in transmitting order



HTTP/2-3  
server

Design 2:

- ALTO specifies that server submits to the HTTP transport in DAG order: submit  $R_i$  only when what  $R_i$  depends on are finished:  $R_1$ ;  $R_2/R_3$ ,  $R_4$
- Issue: Sliding window is large and transport can fit  $R_1/R_2/R_3/R_4$  into a single window



Design 3:

- ALTO indicates the dependencies to HTTP
- Issue: HTTP client can indicate a parent in req header, but this is signaling from client to server; what we need are (1) app signaling to HTTP server, (2) multiple dependencies

# Discuss 3: How/Whether to Specify Settings

- Current draft allows client to specify two (HTTP/2) control knobs on server behaviors
  - 0x02 SETTINGS\_ENABLE\_PUSH (a BCP14 “MUST”)
  - 0x03 SETTINGS\_MAX\_CONCURRENT\_STREAMS (a BCP14 “must”)
- HTTP/3 changed these settings (RFC9114) [SD review]
  - SETTINGS\_ENABLE\_PUSH (0x02): This is removed in favor of the [MAX PUSH ID](#) frame, which provides a more granular control over server push. Specifying a setting with the identifier 0x02 is HTTP/3 error.
  - SETTINGS\_MAX\_CONCURRENT\_STREAMS (0x03): QUIC controls the largest open stream ID as part of its flow-control logic. Specifying it is HTTP/3 error
- Discussion: (1) remove them in the spec and discuss them in operations; (2) specify generic requirements statement

# Discuss 4: Examples using HTTP/1.1 or Later

- Initial draft used HTTP/1.1; then switched to 1.1+2; current draft is only HTTP/2
- Style guide recommends using HTTP/1.1 but if we want to specify more control, we need a way to specify them

```
POST /tqs HTTP/2
Host: alto.example.com
Accept: application/alto-transport+json
Content-Type: application/alto-updatesreamparams+json
Content-Length: TBD

{
  "resource-id": "my-routingcost-map"
}

HTTP/2 200 OK
Content-Type: application/alto-transport+json

{"mq": "/updatesstreams/2718281828459"}
```

Server -> client PUSH\_PROMISE in current stream:

PUSH\_PROMISE

- END\_STREAM

Promised Stream 4

HEADER BLOCK

:method = GET

:scheme = https

:path = /tqs/2718281828459/uq/101

host = alto.example.com

accept = application/alto-error+json,  
application/alto-costmap+json

Server -> client content Stream 4:

HEADERS

+ END\_STREAM

+ END\_HEADERS

:status = 200

content-type = application/alto-costmap+json

content-length = TBD

DATA

+ END\_STREAM

{

"meta" : {

"dependent-vtags" : [ {

"resource-id": "my-network-map",

"tag": "da65eca2eb7a10ce8b059740b0b2e3f8eb1d4785"

# Discuss 5: Media Type

- Finalize media type: a single media type (`application/alto-transport+json`) to encode all new data items

# Next Step

- The authors will submit a new version as soon as the WG makes the decisions on Discuss 1-5 (in two weeks if possible)
  - Discuss 1.1 Allow client long pull
  - Discuss 1.2 Server push promise -> server put | + server put
  - Discuss 2: How much to specify ordering control: how much transport-awareness of app semantics
  - Discuss 3: How/whether to specify settings
  - Discuss 4: Examples using HTTP/1.1 or HTTP/2
  - Discuss 5: Media type comments

# Backup Slides

Service	Type (Out if not labeled)	Core Information Structures	Main Size Var	Stability Expectation	Incremental Changes
Information Resource Directory [RFC7285]	application/alto-directory+json	Key-value store; Delegation	#resources	Stable	Add/delete resources
Network Map [RFC7285]	application/alto-networkmap+json	Key-value store: pid -> addrType -> array	#CIDRs	Stable	Add/delete CIDR from pid
Cost Map [RFC7285]	application/alto-costmap+json	Key-value store: srcPID -> dstPID -> value Depend on network map	#SRCID * #DSTPID	More dynamic than network map	Update cost map entries
Filtered Map Services [RFC7285]	In: application/alto-networkmapfilter+json Out: application/alto-costmapfilter+json	In: selected srcPID, dstPID Key-value store: srcPID -> dstPID -> value Depend on network map	Filtered #SRCID * #DSTPID	More dynamic than network map	Update cost map entries
Endpoint Property Service [RFC7285]	In: application/alto-endpointpropparams+json Out: application/alto-endpointprop+json	In: addr + prop Key-value store: addr -> prop -> value	#addr * #prop	Depend on property, can be dynamic or stable	Update property value
Endpoint Cost Service [RFC7285]	In: application/alto-endpointcostparams+json Out: application/alto-endpointcost+json	In: srcAddr x dstAddr Key-value store: srcAddr -> dstAddr -> value	#src * #dst	Can be more dynamic than cost map	Update cost value
Cost Calendar [RFC8896]	Extension to CostType	Calendar array	Previous * #num_intervals	Can be dynamic	Calendar window moves
Unified Property	In: application/alto-propmapparams+json; Out: application/alto-propmap+json	In: addr, prop Key-value store: addr -> prop -> value	#addr * #prop	Depend on property, can be dynamic or stable	Update property value
Path Vector	In: see costmap; Out: multipart/related;type=application/alto-costmap+json,	Cost map + unified property map	#src * #dst * #vec size	Depend on metric, can be dynamic or stable	Update path vector
CDNi Cap & Footprint	application/alto-cdni+json	Array; {capability-type: capability-value}	#footprint * #capability	Can be dynamic, bursty	Update capabilities

# Performance and Effectiveness of Current Transport

Infrastructures	Basic Workload (ALTO SPEC)	Transport	Collecting metrics
<ul style="list-style-type: none"><li>• Benocs is fully open to use its infrastructure as an evaluation environment</li><li>• Greater Bay Network is also fully open to use its infrastructure as an evaluation environment</li></ul>	<ol style="list-style-type: none"><li>1. (Filtered) Cost map: distribute inter-site performance metrics and calendar; routing changes + link dynamics =&gt; updated metrics</li><li>2. Endpoint/unified property service: endpoint access status query/updates/bwe</li><li>3. CDN node footprint &amp; capability</li><li>4. Flow direction (pointing to CDN nodes) using ECS</li><li>5. Path vector providing available reservable bandwidth</li></ol>	<ul style="list-style-type: none"><li>• HTTP/1.x per request full retrieval<ul style="list-style-type: none"><li>• keep alive</li><li>• pipelining</li></ul></li><li>• HTTP/2, HTTP/3 per request, full retrieval</li><li>• ALTO/SSE RFC8895 (on HTTP/1.x)</li></ul>	<ul style="list-style-type: none"><li>• ALTO Server processing load</li><li>• ALTO client processing load</li><li>• Transport load (bytes)</li><li>• Transport latency</li><li>• Throughput</li><li>• Scalability</li></ul>

# General Space of Network->App Information Transport

Protocol/Base Reference	What information	How transported	Frequency
ALTO/RFC7285, RFC8895	Network map, cost map, unified property, ...	HTTP/1.x client/server	Request/response; ALTO/SSE incremental push
ECN/RFC3168	Congestion notification	2 bits in IP Traffic Class header; ECN-Echo/CWR flags in TCP header	Per packet marking
NEF, SCEF<->AS/3GPP, ...	Network capabilities/events -> AF	HTTP	Request/response
...			

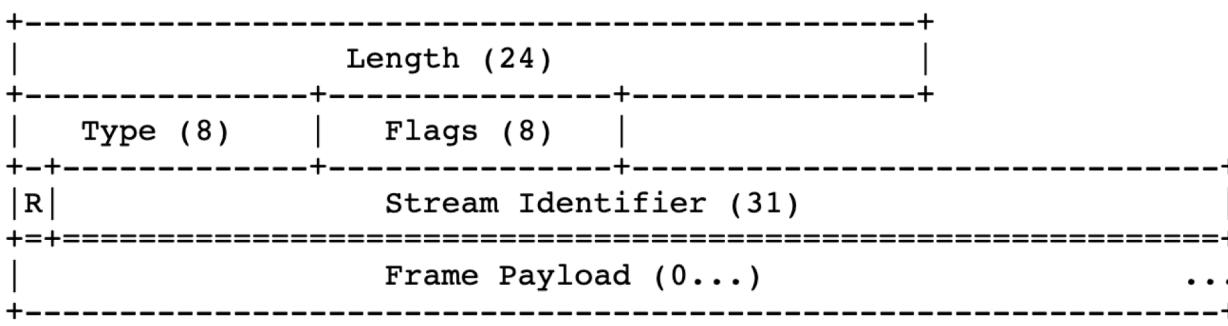


Figure 1: Frame Layout

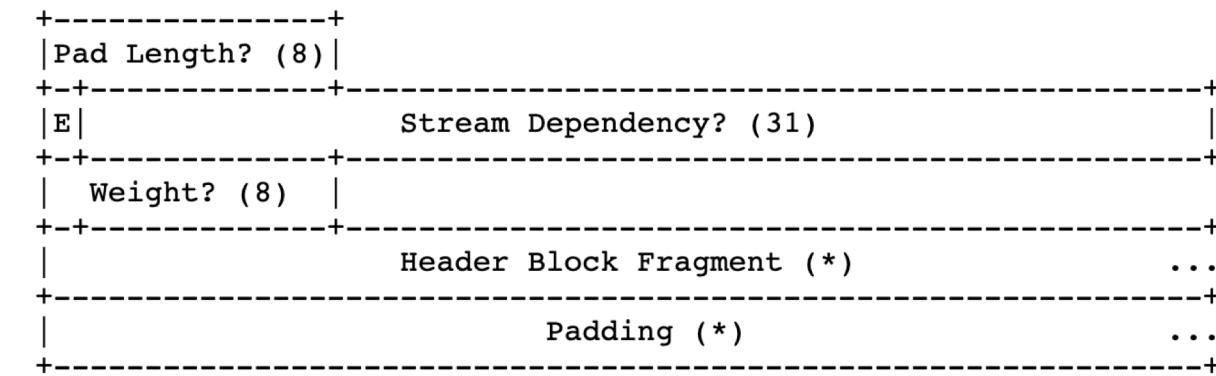


Figure 7: HEADERS Frame Payload

The value 0x0 is reserved for frames that are associated with the connection as a whole as opposed to an individual stream.

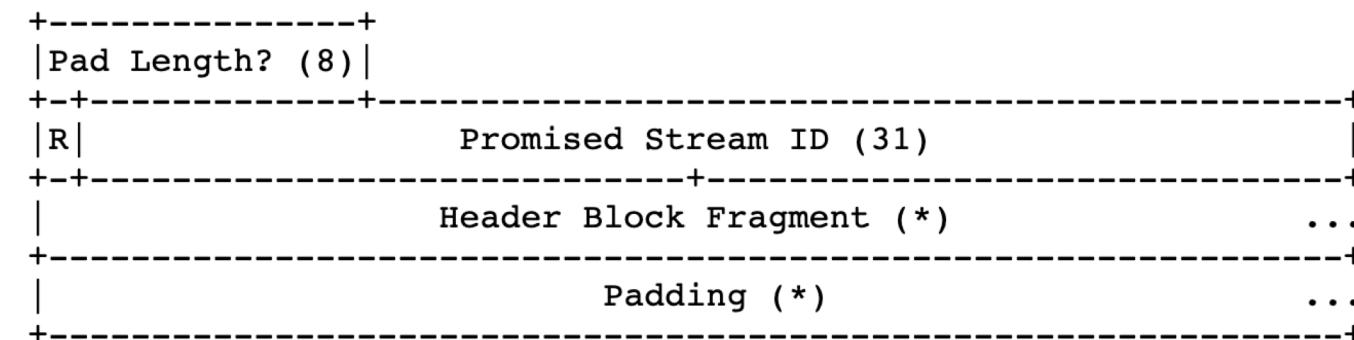


Figure 11: PUSH\_PROMISE Payload Format

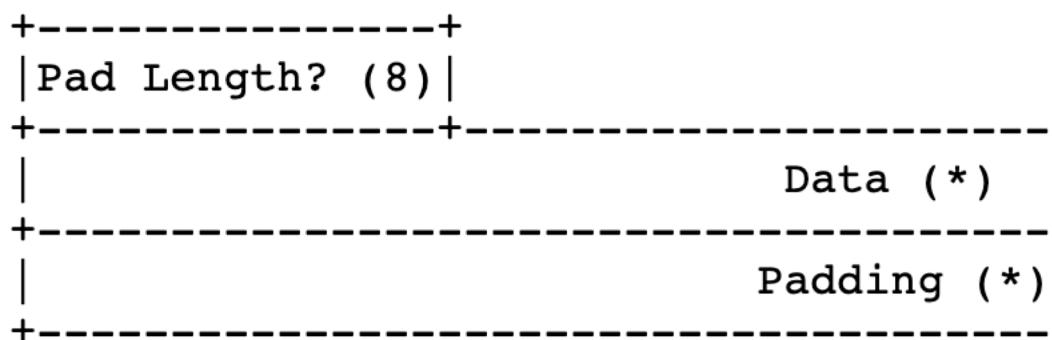


Figure 6: DATA Frame Payload

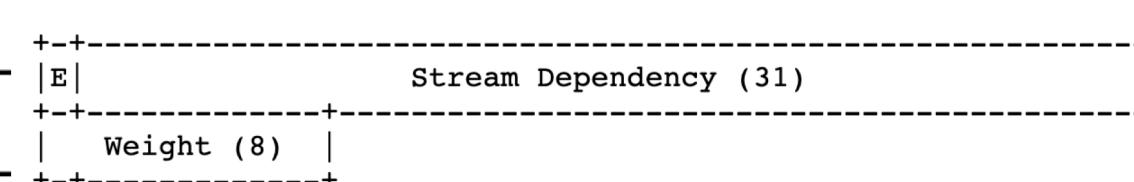
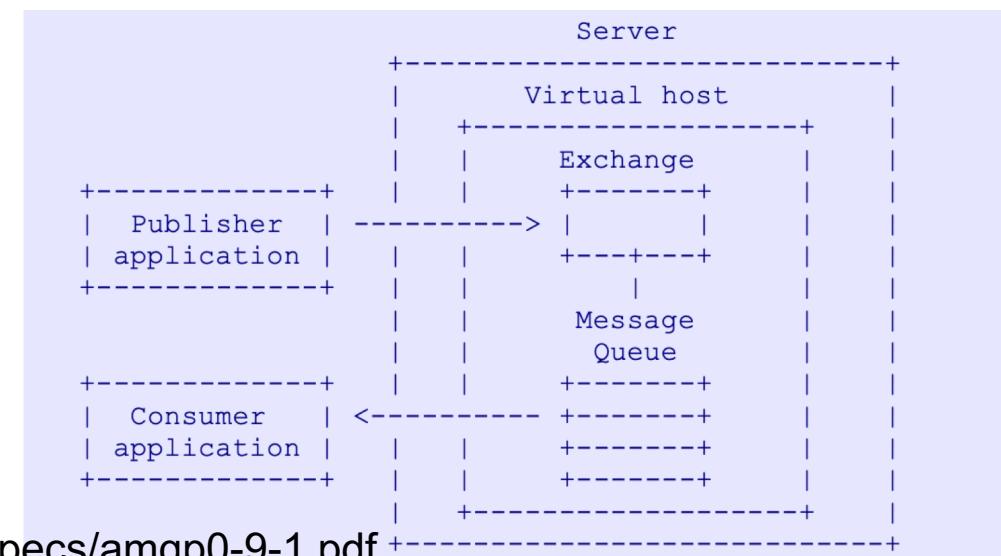


Figure 8: PRIORITY Frame Payload

# Transport and Pub/sub

- What is missing
  - The design does not allow creation of generic message queues
  - Only the server can be the publisher
    - Clients publish info to be shared with other client
  - The design does not have the capability of Exchange (message router)
- Way forward: Keep simple
  - Broker for further discussion



<https://www.rabbitmq.com/resources/specs/amqp0-9-1.pdf>

# Additional Information about Transport Queue

- Calendar semantics
  - Tell the client ALTO information (e.g., cost) for a future time point
  - Tell the client when the next information will be released, it is the time that the info is released is distributed, not the value [support]

# Negotiation

- Task: to fully specify the complete set of negotiation parameters
  - Use HTTP/2 (default), HTTP/3
  - Incremental updates queue encoding
    - Chosen by server, but must be specified as HTTP Accept header of client
    - Initial connection setup: client sends list of incr update mime types, and server returns the subset which it can use (Vary header in response)
  - Concurrency level (e.g., Rucio controller need to monitor info for a large number of clients) stream control